
MLDS HW3-1

TAs

b05902013@csie.ntu.edu.tw

b05902127@csie.ntu.edu.tw

Outline

- ❖ **Timeline**
- ❖ **Task Descriptions**
- ❖ **Model & Training tips**
- ❖ **Submission & Rules**
- ❖ **Q&A**

Timeline

Three Parts in HW3

- (3-1) Image Generation
- (3-2) Text-to-Image Generation
- (3-3) Style Transfer


Schedule

- **4/30 or 5/4 :**
 - Release HW3-1
- **5/7 or 5/11 :**
 - Present HW 3-1
 - Release HW 3-2, HW 3-3
- **5/14 or 5/18 : Break**
- **5/21 or 5/25:**
 - Present HW 3-2, HW 3-3


Task Descriptions

HW3-1: Image Generation _{2/2}

**Bird
Generative Model**




**Flower
Generative Model**



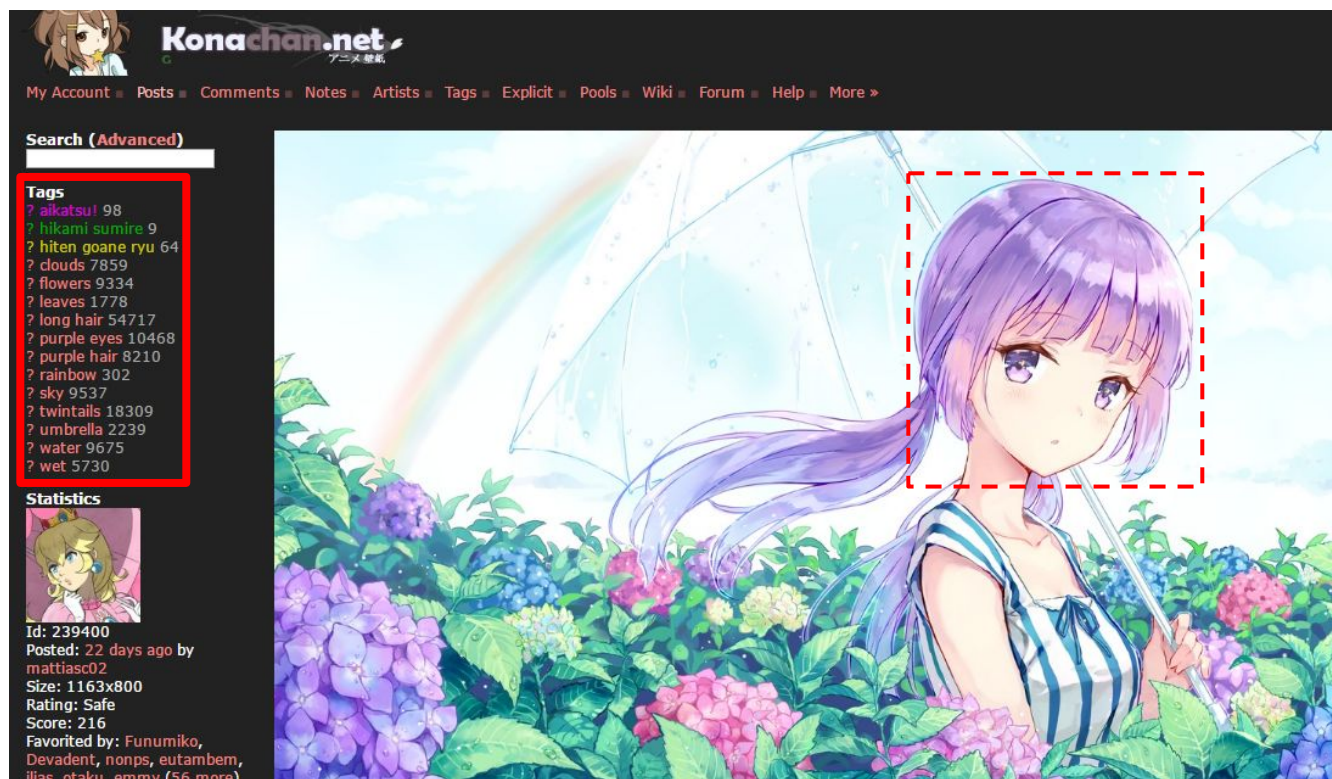
HW3-1: Image Generation _{2/2}

**Anime
Generative Model**



Data Collections ^{1/2}

- Anime dataset




http://konachan.net/post/show/239400/aikatsu-clouds-flowers-hikami_sumire-hiten_goane_r

感謝樊恩宇助教蒐集data

Data Collections ^{2/2}

- Extra data

MakeGirlsMoe Home History Transition Help ▾



Generate

👍 +1

👎 -1

Share on Twitter

Options

☐ Advanced Mode

Model

Camellia 256x256 Ver.171219 (9.9MB)

Hair Color

Random

Hair Style

Random

Eye Color

Random

Dark Skin

Off

Random

On

Blush

Off

Random

On

Smile

Off

Random

On

Open Mouth

Off

Random

On

Hat

Off

Random

On

Ribbon

Off

Random

On

Glasses

Off

Random

On

Style

Random

Noise

Random

Fixed

Current Noise

Noise Import/Export


Import

Export

Model & Training Tips

GAN ^{1/5}

- Overview



GAN 2/5

- Discriminator

- input: images (batch_size, height, width, channels)
- output: scores (batch_size,)
- architecture: CNN, DNN

- Generator

- input: noises (batch_size, noise_dim)
- output: images (batch_size, height, width, channels)
- architecture: CNN, DNN
- use deconvolution (transpose convolution) layer in CNN

Baseline Model

- Generator

- input = (100,)
- Dense(128*16*16, 'relu')
- Reshape((16, 16, 128))
- Upsampling
- Conv2D(128, kernel = 4)
- Relu
- Upsampling
- Conv2D(64, kernel = 4)
- Relu
- Conv2D(3, kernel = 4)
- tanh

- Training

- Adam(lr = 0.0002, beta = 0.5)

- Discriminator

- input = (64, 64, 3)
- Conv2D(32, kernel = 4)
- Relu
- Conv2D(64, kernel = 4)
- ZeroPadding
- Relu
- Conv2D(128, kernel = 4)
- Relu
- Conv2D(256, kernel = 4)
- Relu
- Flatten
- Dense(1, sigmoid)

GAN ^{3/5}

- Training procedure
 - repeat max_iteration times
 - repeat d_update times
 - discriminator = **update_discriminator**(training_data, generator)
 - repeat g_update times
 - generator = **update_generator**(discriminator)
- Testing procedure
 - noise = sample_batch_noise(batch_size, noise_dim)
 - output_images = generator(noise)

GAN 4/5

- Update discriminator

- `real_images = sample_batch_data(training_data, batch_size)`
- `noise = sample_batch_noise(batch_size, noise_dim)`
- `fake_images = generator(noise)`
- `real_predicts = discriminator(real_images)`
- `fake_predicts = discriminator(fake_images)`
- `d_loss = loss_d_fn(real_predicts, real_labels, fake_predicts, fake_labels)`
- `d_grad = gradients(d_loss, d_params)`
- `d_params = updates(d_params, d_grad)`
- `# do not` update the parameters of generator

- Update Generator

- `noise = sample_batch_noise(noise_dim, batch_size)`
- `fake_images = generator(noise)`
- `fake_predicts = discriminator(fake_images)`
- `g_loss = loss_g_fn(fake_predicts, real_labels)`
- `g_grad = gradients(g_loss, g_params)`
- `g_params = updates(g_params, g_grad)`
- `# do not` update the parameters of discriminator

Wasserstein GAN ^{1/3}

The output of D is thus not probability anymore.
The D loss turn to be a measure of distance.

$$L_D^{WGAN} = E[D(x)] - E[D(G(z))]$$

$$L_G^{WGAN} = E[D(G(z))]$$

$$W_D \leftarrow \text{clip_by_value}(W_D, -0.01, 0.01)$$

ref:<https://arxiv.org/abs/1701.07875>

Wasserstein GAN 2/3

- In each training iteration: **No sigmoid for the output of D**

Learning
D

Repeat
k times

- Sample m examples $\{x^1, x^2, \dots, x^m\}$ from data distribution $P_{data}(x)$
- Sample m noise samples $\{z^1, z^2, \dots, z^m\}$ from the prior $P_{prior}(z)$
- Obtaining generated data $\{\tilde{x}^1, \tilde{x}^2, \dots, \tilde{x}^m\}$, $\tilde{x}^i = G(z^i)$
- Update discriminator parameters θ_d to maximize
 - $\tilde{V} = \frac{1}{m} \sum_{i=1}^m D(x^i) - \frac{1}{m} \sum_{i=1}^m D(\tilde{x}^i)$
 - $\theta_d \leftarrow \theta_d + \eta \nabla \tilde{V}(\theta_d)$ **Weight clipping**

Learning
G

Only
Once

- Sample another m noise samples $\{z^1, z^2, \dots, z^m\}$ from the prior $P_{prior}(z)$
- Update generator parameters θ_g to minimize
 - $\tilde{V} = \frac{1}{m} \sum_{i=1}^m \log D(x^i) - \frac{1}{m} \sum_{i=1}^m D(G(z^i))$
 - $\theta_g \leftarrow \theta_g - \eta \nabla \tilde{V}(\theta_g)$

Wasserstein GAN ^{3/3}

- Implementation Notes:
 - Do not apply sigmoid at the output of D
 - Clip the weight of D
 - Use RMSProp instead of Adam
 - Train more iteration of D (the paper use 5)

Improved WGAN (WGAN-GP) ^{1/2}

Do not clip the weight of D but to add a new objective called “Gradient Penalty”.

$$L_D^{WGAN_GP} = L_D^{WGAN} + \lambda E[(\|\nabla D(\alpha x + (1 - \alpha)G(z))\|_2 - 1)^2]$$


$$L_G^{WGAN_GP} = L_G^{WGAN}$$

ref:<https://arxiv.org/pdf/1704.00028.pdf>

Improved WGAN (WGAN-GP) ^{2/2}

$$W(P_{data}, P_G) \approx \max_D \{ E_{x \sim P_{data}} [D(x)] - E_{x \sim P_G} [D(x)] - \lambda E_{x \sim P_{penalty}} [\max(0, \|\nabla_x D(x)\| - 1)] \}$$

$$(\|\nabla_x D(x)\| - 1)^2$$



ref:<https://arxiv.org/pdf/1704.00028.pdf>

Least Squares GAN

$$\begin{aligned}
 \min_D V_{\text{LSGAN}}(D) &= \frac{1}{2} \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [(D(\mathbf{x}) - b)^2] + \frac{1}{2} \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [(D(G(\mathbf{z})) - a)^2] \\
 \min_G V_{\text{LSGAN}}(G) &= \frac{1}{2} \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [(D(G(\mathbf{z})) - c)^2],
 \end{aligned}
 \tag{2}$$

ref:<https://arxiv.org/abs/1611.04076>

Little Results



GAN result



WGAN result

TA Results



DCGAN



DCGAN w/ GP


dataset : <https://crypko.ai/#/>

Submission & Grading

Data & format

- Anime Dataset

- training data: 33.4k (image, tags) pair
- **faces/**, tags.csv, sample_testing_text.txt



blue eyes
red hair
short hair

- training tags file format

- img_id <comma> tag1 <colon> #_post <tab> tag2 <colon> #_post
- tags.csv is **not** used in HW3-1

```
1 0,touhou:17705 |chen:423 |moneti daifuku :60 |animal ears:12241 |catgirl:4903 |
2 1,touhou:17697 |onozuka komachi:224 |shikieiki yamaxanadu:217 |$
3 2,original:25774 |blonde hair:25457 |doll:1040 |dress:16585 |pink eyes:3896 |ta
4 3,amagi brilliant park:111 |musaigen no phantom world:39 |nichijou:142 |kawakan
```

tags.csv

- testing text file format

- testing_text_id <comma> testing_text
- testing text only includes '**color hair**' and '**color eyes**', only alphabetic char involved.
- sample_testing_text.txt is **not** used in HW3-1

```
1 1,blue hair blue eyes
2 2,blue hair green eyes
3 3,blue hair red eyes
4 4,green hair blue eyes
```

sample
testing_text.txt

Data & format

- Extra data

- training data: 36.7k (image, tags) pair
- **images/**, tags.csv



black eyes
red hair

- training tags file format

- img_id <comma> hair tag <space> eyes tag
- tags in extra data only includes 'color hair' and 'color eyes'
- tags.csv is **not** used in HW3-1

```
1 0,aqua hair aqua eyes
2 1,aqua hair aqua eyes
3 2,aqua hair aqua eyes
4 3,aqua hair aqua eyes
5 4,aqua hair aqua eyes
```

tags.csv

Data Link

- [Anime Dataset](#)
- [Extra Data](#)
- [Extra Data2 \(no tag\)](#)

HW3 Grading Policy:

- HW3-1 Code (image generation) : 5%
- HW3-2 Code (text-to-image generation): 5%
- Report : 15%
- HW3-3 (Bonus, style transfer): 2%
- 分工表 : 0.5%

HW3-1 Report Questions

- Model Description
 - Describe the models you use to, including the model architecture, objective function for G and D. (1%)
- Experiment settings and observation
 - Show generated images (1%)
- Compare your model with WGAN, WGAN-GP, LSGAN (choose 1)
 - Model Description of the choosed model (1%)
 - Result of the model (1%)
 - Comparison Analysis (1%)
- Training tips for improvement (4%)
- Mode Collapse (2%)

Training Tips for improvement

- Pick **three** tips in the following website
 - <https://github.com/soumith/ganhacks>
 - Please implement these tips on image generation
- Total : 3%, 1% for each
 - Which tip & implement details (0.5%)
 - Result (image or loss...etc.) and Analysis (0.5%)
- Only the following tips are accepted
 - 1, 2, 3, 4, 5, 6, 9, 13, 14, 17

Mode Collapse

- First, train a model that leads to mode collapse and record its model architecture as well as hyperparameters (especially iterations and lr)
- Then, find ways to alleviate mode collapse phenomenon without modifying model architecture, lr and iterations

HW3-1 Code Grading

- Reproduce Score : 2%
- Baseline Score : 2%
- TA Review : 1% (**mode collapse**)



Output Format Requirement

- The generated images should be in Directory **samples/**
 - 請大家繳交時**就將產生的結果傳到samples/**, 助教利用script reproduce時也請同學將結果輸出到這個資料夾。為保證reproduce結果相同, 請同學將random的部份固定
 - 批改作業會在azure上, 請同學繳交前在機台上檢查
 - 已經在github裡的image -> samples/gan_original.png
 - run_gan.sh -> samples/gan.png
- Each generated image must be resized to **64 x 64**
- Generate 25 image into one png
 - sample code is in baseline.py
 - 為防止同學產生的圖片不一致, 請同學使用baseline.py裡的**save_imgs()**

Baseline Model

- [Anime Face Recognition](#)
- [Github for baseline.py](#)
- How to run:
 - pip install opencv-python
 - Download pre-trained model in [Here](#).
 - python baseline.py --input <input_image>
- Generate 25 images in one png
 - if faces > 20: pass

Recommended Packages

- Python 3.6
- **Tensorflow**
- **PyTorch**
- Keras
- MXNet
- matplotlib
- skimage
- Python Standard Library

Submission on Github

- Only one branch **master** is needed
- Only **generator** and **inference** mode is needed
- Remember to put your **pre-trained models or download scripts** so that we can run your code successfully

Q&A

b05902013@csie.ntu.edu.tw

b05902127@csie.ntu.edu.tw