

---

# MLDS 2019 Spring HW4-2 - Deep Q Learning

2019/06/01  
adlxmls@gmail.com

---

---

# Time Schedule

- June 11th 19:30 Deadline
- June 15th 10:30 Deadline

# Outline

	HW4-1	HW4-2	HW4-3
Pong	PG		AC
Breakout		DQN	AC
Improved Version			

## Outline

- Introduction
  - Game Playing : Breakout
- Deep Reinforcement Learning
  - Deep Q-Learning (DQN)
  - Improvements to DQN
- Grading & Format
  - Grading Policy
  - Code Format
  - Report
  - Submission

# Introduction

## Environment

### Breakout




# Deep Q-Learning (DQN)

	a1	a2
s1	-3	1
s2	-1	3
s3	3	3
s4	2	-2
...		

# Deep Q-Learning (DQN)

	a1	a2
s1	-3	1
s2	-1	3
s3	3	37
s4	2	-20
...		




# Deep Reinforcement Learning

## Deep Q-Learning (DQN)

“classic” deep Q-learning algorithm:

Replay buffer

- 
1. take some action  $\mathbf{a}_i$  and observe  $(\mathbf{s}_i, \mathbf{a}_i, \mathbf{s}'_i, r_i)$ , add it to  $\mathcal{B}$
  2. sample mini-batch  $\{\mathbf{s}_j, \mathbf{a}_j, \mathbf{s}'_j, r_j\}$  from  $\mathcal{B}$  uniformly
  3. compute  $y_j = r_j + \gamma \max_{\mathbf{a}'_j} Q_{\phi'}(\mathbf{s}'_j, \mathbf{a}'_j)$  using *target* network  $Q_{\phi'}$
  4.  $\phi \leftarrow \phi - \alpha \sum_j \frac{dQ_\phi}{d\phi}(\mathbf{s}_j, \mathbf{a}_j)(Q_\phi(\mathbf{s}_j, \mathbf{a}_j) - y_j)$
  5. update  $\phi'$ : copy  $\phi$  every  $N$  steps



# Deep Reinforcement Learning

## Deep Q-Learning (DQN)

- The action should act  $\epsilon$ -greedily
  - Random action with probability  $\epsilon$
- Linearly decline  $\epsilon$  from 1.0 to some small value, say 0.025
  - Decline per step
  - Randomness is for exploration, agent is weak at start
- Hyperparameters
  - Replay Memory Size 10000
  - Perform Update Current Network Step 4
  - Perform Update Target Network Step 1000
  - Learning Rate  $1.5e-4$
  - Batch Size 32

# Deep Reinforcement Learning


## Improvements to DQN

- Double Q-Learning
- Dueling Network
- Prioritized Replay Memory
- Noisy DQN
- Distributional DQN

<https://arxiv.org/pdf/1710.02298.pdf>

# Training Tips

## Training Plot



- X-axis : 1000 episodes/unit
- Y-axis : **Unclipped** reward per episode

# Training Tips

## Why Reward is clipped

- Performing the same action for 4 frames
  - To use data more efficiently
- Reward may be up to 4
  - If positive, clip to 1 → reduce variance
- How to see your unclipped reward
  1. Use the *test* function
  2. Turn off the *clip\_reward* option of your environment and do the clipping by yourself.

# Training Tips

## Asynchronous Update (Optional)

- In tensorflow, *feed\_dict* does the copy thing
  - Upon updating, the agent have to wait for it to continue exploring.
- Try run the update asynchronously
  - Main thread : Collect data
  - The other thread : Copy data to GPU
  - GPU : Training
  - Using the thread/multiprocessing module
- This is totally **not necessary for you to get baseline**, just some speed-up you can try.
  - This can go wrong and annoying if you' re not familiar with threading, thus I recommend not to try it unless you are confident enough.

## Code Format

- Please download the sample files from [github](#)
- Follow the instructions in README to install required packages
- **Four** functions you should implement in [agent\\_dqn.py](#)
  1. `__init__(self, env, args)`
  2. `init_game_setting(self)`
  3. `train(self)`
  4. `make_action(self, state, test)`
- **DO NOT** add any parameter in `__init__()`, `init_game_setting()` and `make_action()`
- You can add new methods in the [agent\\_dqn.py](#)
- You can add your arguments in [argument.py](#)

## Submission

- Submit your presentation files to: [google drive](#)

## Baseline

- DQN
  - Getting averaging reward in 100 episodes over 40 in **Breakout**
  - With **OpenAI's Atari wrapper** & reward clipping
    - You SHOULD will unclip the reward when testing



## Slides

- Describe your DQN model
- Plot the learning curve to show the performance of your Deep Q Learning on Breakout
  - X-axis: number of time steps
  - Y-axis: average reward in last 30 episodes
- Implement **1** improvement method on page 10
  - Describe your tips for improvement
  - Learning curve
  - Compare to origin Deep Q Learning

# Related Materials

- Course & Tutorial:
  - [Berkeley Deep Reinforcement Learning, Fall 2017](#)
  - [David Silver RL course](#)
  - [Nips 2016 RL tutorial](#)
- Blog:
  - [Andrej Karpathy's blog](#)
  - [Arthur Juliani's Blog](#)
- Text Book:
  - [Reinforcement Learning: An Introduction](#)
- Repo:
  - <https://github.com/williamFalcon/DeepRLHacks>

# Deep Reinforcement Learning

## Double DQN

- The formula  $Y_t^{\text{DQN}} \equiv R_{t+1} + \gamma \max_a Q(S_{t+1}, a; \theta_t^-) . \quad (3)$  often overestimates the maximum Q value.
- Thus instead choose the action of the max Q in the target network, choose the action of the max Q in the **current network**.

$$Y_t^{\text{DoubleDQN}} \equiv R_{t+1} + \gamma Q(S_{t+1}, \underset{a}{\operatorname{argmax}} Q(S_{t+1}, a; \theta_t), \theta_t^-) .$$

# Deep Reinforcement Learning

## Dueling Network

- In many state, action does not counts.
  - DQN tries to find out the max  $Q$  in each state
- Use same network to output *Value* and *Advantage*

- Why should it be the *Advantage*?

- Add loss constraint


$$Q(s, a; \theta, \alpha, \beta) = V(s; \theta, \beta) +$$

$$\left( A(s, a; \theta, \alpha) - \max_{a' \in |\mathcal{A}|} A(s, a'; \theta, \alpha) \right)$$

- Alternative  $Q$  function, more stable (more used)

$$Q(s, a; \theta, \alpha, \beta) = V(s; \theta, \beta) +$$

$$\left( A(s, a; \theta, \alpha) - \frac{1}{|\mathcal{A}|} \sum_{a'} A(s, a'; \theta, \alpha) \right)$$



# Deep Reinforcement Learning

## Prioritized Replay Memory

- DQN : Sample from replay memory uniformly
  - We can sample the replays with large loss more often
  - Thus we sample with the probability
    - TD ERROR =  $R_j + \gamma_j Q_{\text{target}}(S_j, \arg \max_a Q(S_j, a)) - Q(S_{j-1}, A_{j-1})$
    - $p_t \propto |TD \text{ ERROR}|^\omega$
    - $\omega$  is a hyperparameter, 0.5 in Rainbow
- $$p_t \propto \left| R_{t+1} + \gamma_{t+1} \max_{a'} q_{\bar{\theta}}(S_{t+1}, a') - q_{\theta}(S_t, A_t) \right|^\omega$$

<https://arxiv.org/pdf/1511.05952.pdf>

# Deep Reinforcement Learning

## Prioritized Replay Memory

- However, the resulting gradient estimator is biased, since we are sampling from a different distribution

- Correct by importance sampling weights

- With  $\rho_i = 1 / P(i)$ , the IS weights

$$w_i = \left( \frac{1}{N} \cdot \frac{1}{P(i)} \right)^\beta$$

- $\beta$  is linearly declined to 1

- $\beta = 1 \rightarrow$  Unbiased


- Try to learn quicker  $\rightarrow$  Try to converge correctly

$$\tilde{v}_g \doteq \frac{\sum_{k=1}^n \rho_k Y_k}{n}.$$

# Deep Reinforcement Learning

## Prioritized Replay Memory

- Using array, the complexity of sampling is  $O(n)$ 
  - Try another data structure
- Sum Tree, which prioritized sampling can be  $O(\lg n)$ 
  - Devide the priorities into k groups(batch size) by the max priority
  - That is if the max is 42, batch size = 6, we devide them into [1, 7], [8, 14],  $\dots$ , [36, 42]
  - Randomly sample a number from each interval
  - Go down the sum tree by the priority to retrieve the data at the leaf



# Deep Reinforcement Learning

## Prioritized Replay Memory

---

**Algorithm 1** Double DQN with proportional prioritization

---

```
1: Input: minibatch  $k$ , step-size  $\eta$ , replay period  $K$  and size  $N$ , exponents  $\alpha$  and  $\beta$ , budget  $T$ .
2: Initialize replay memory  $\mathcal{H} = \emptyset$ ,  $\Delta = 0$ ,  $p_1 = 1$ 
3: Observe  $S_0$  and choose  $A_0 \sim \pi_\theta(S_0)$ 
4: for  $t = 1$  to  $T$  do
5:   Observe  $S_t, R_t, \gamma_t$ 
6:   Store transition  $(S_{t-1}, A_{t-1}, R_t, \gamma_t, S_t)$  in  $\mathcal{H}$  with maximal priority  $p_t = \max_{i < t} p_i$ 
7:   if  $t \equiv 0 \pmod K$  then
8:     for  $j = 1$  to  $k$  do
9:       Sample transition  $j \sim P(j) = p_j^\alpha / \sum_i p_i^\alpha$ 
10:      Compute importance-sampling weight  $w_j = (N \cdot P(j))^{-\beta} / \max_i w_i$ 
11:      Compute TD-error  $\delta_j = R_j + \gamma_j Q_{\text{target}}(S_j, \arg \max_a Q(S_j, a)) - Q(S_{j-1}, A_{j-1})$ 
12:      Update transition priority  $p_j \leftarrow |\delta_j|$ 
13:      Accumulate weight-change  $\Delta \leftarrow \Delta + w_j \cdot \delta_j \cdot \nabla_\theta Q(S_{j-1}, A_{j-1})$ 
14:    end for
15:    Update weights  $\theta \leftarrow \theta + \eta \cdot \Delta$ , reset  $\Delta = 0$ 
16:    From time to time copy weights into target network  $\theta_{\text{target}} \leftarrow \theta$ 
17:  end if
18:  Choose action  $A_t \sim \pi_\theta(S_t)$ 
19: end for
```

---